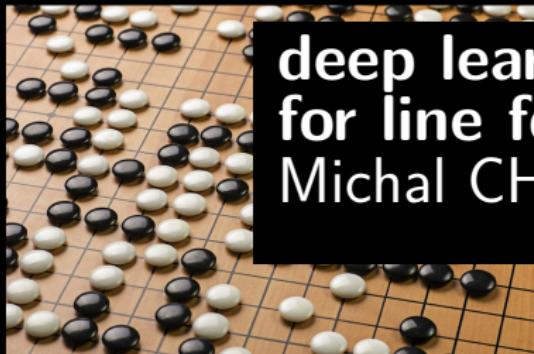


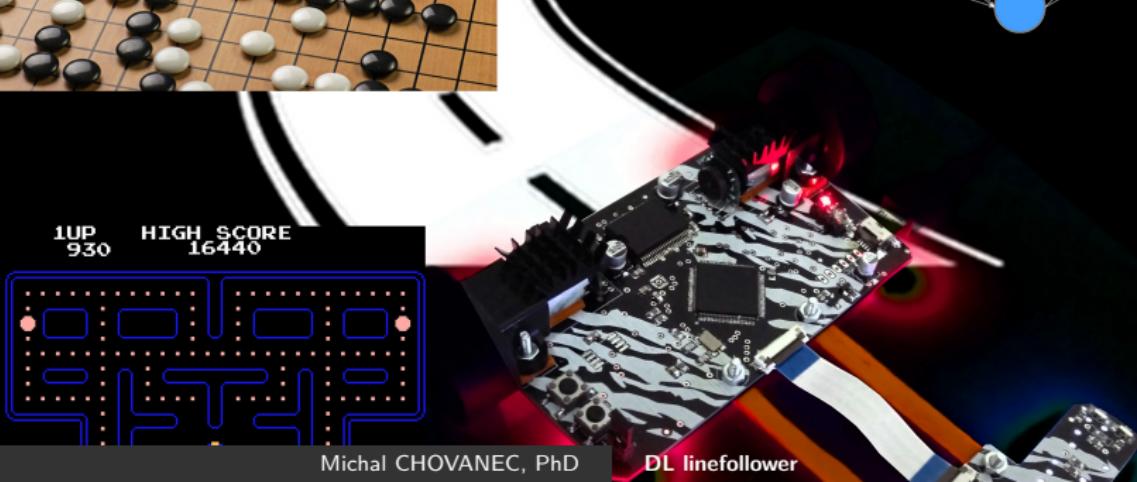
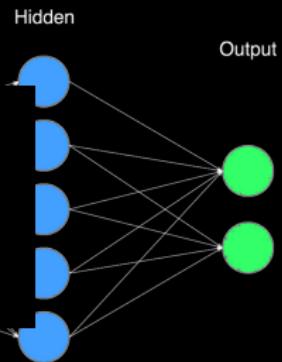
$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [r_{t+1} + \lambda \max_a Q(s_{t+1}, a) - Q(s_t, a_t)]$$

(The New Action Value = The Old Value) + The Learning Rate \times (The New Information — the Old Information)



deep learning for line following robot

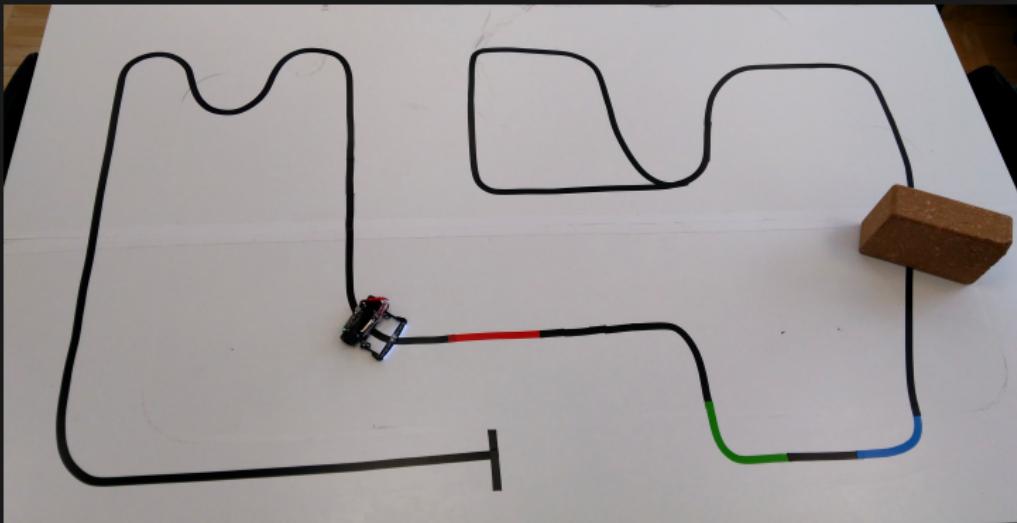
Michal CHOVANEC



Michal CHOVANEC, PhD

DL linefollower

line follower competition



- SR - Istrobot
- CR - Roboticky den
- AU - Robot Challenge
- PL - Zawody robotow

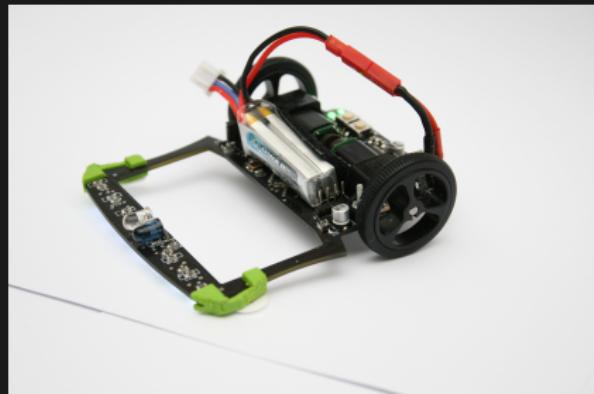
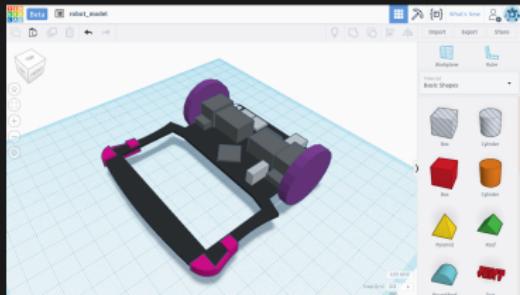
what does it take

- Hardware

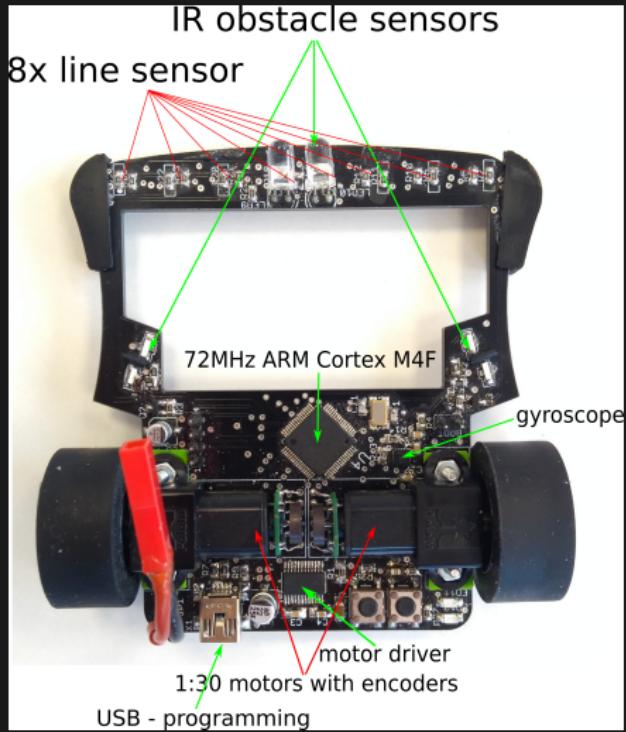
- strong light motors
- high adhesion tyres
- light accumulator
- fast CPU
- dozen of sensors

- Software

- hard real time OS
- tunned PIDs
- predictive controll
- mapping

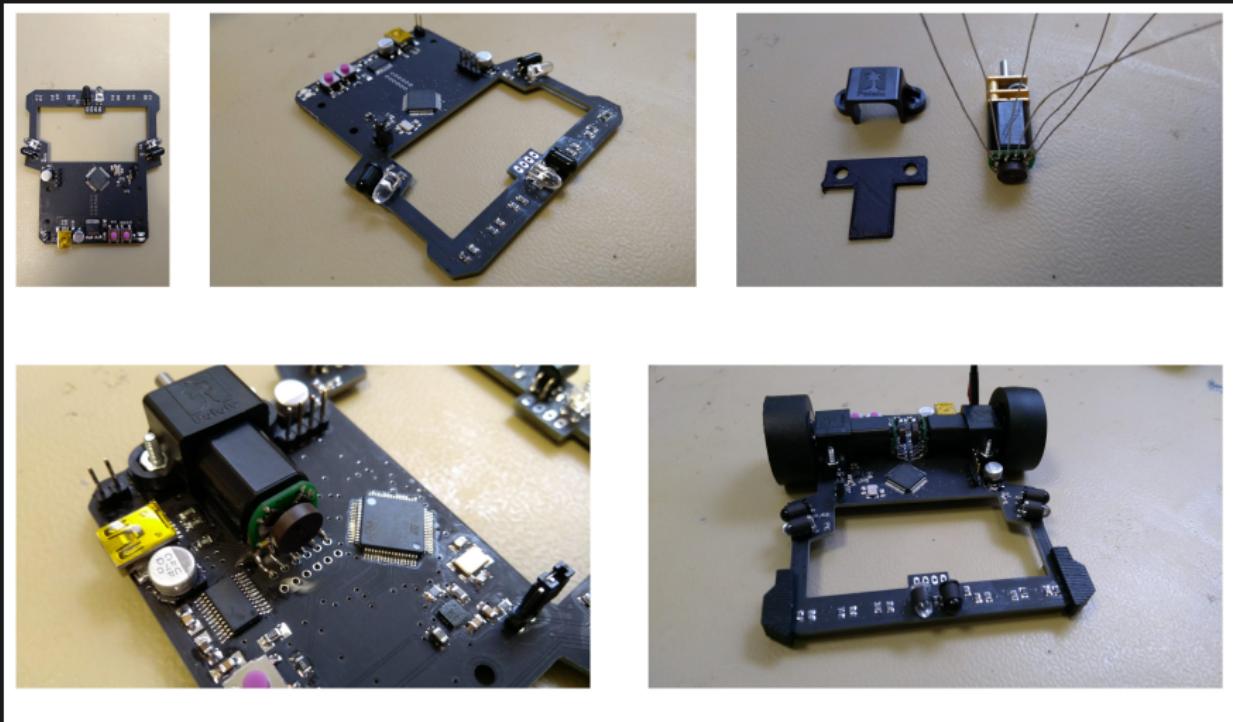


motoko uprising - hardware



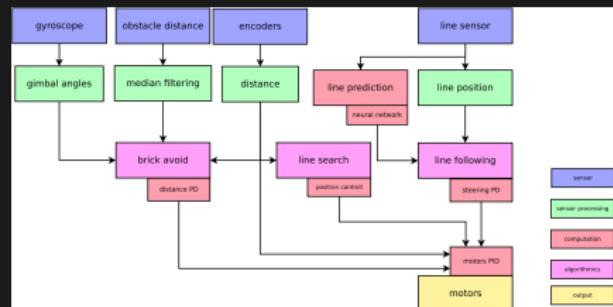
- **mcu** : STM32F303, 72MHz ARM Cortex M4F
- **motor driver** : TI DRV8834
- **motors** : 1:30 HP Pololu, with magnetic encoder
- **tyres** : Pololu 28mm diameter
- **line sensor** : 8x 540nm phototransistor + white LED
- **obstacle sensor** : SMD IR phototransistor + IR LED
- **imu** : LSM6DS0, gyroscope + accelerometer

motoko uprising - hardware



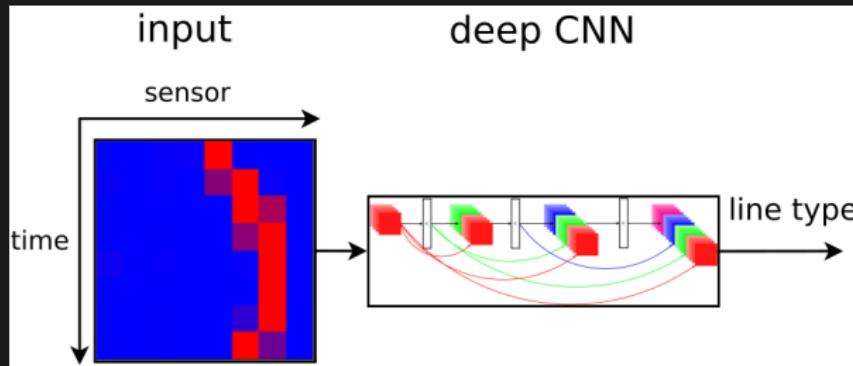
software

- **quadratic interpolation** : high precision line position computing
- **steering PID** : PD controller for steering
- **motor PID** : two PIDs for motors speed controll
- **curve shape prediction**
 - **fast** run on straight line, **brake** on curve
 - **deep neural network**
- written in **C++**
- **network training** on GPU - own framework for CNN

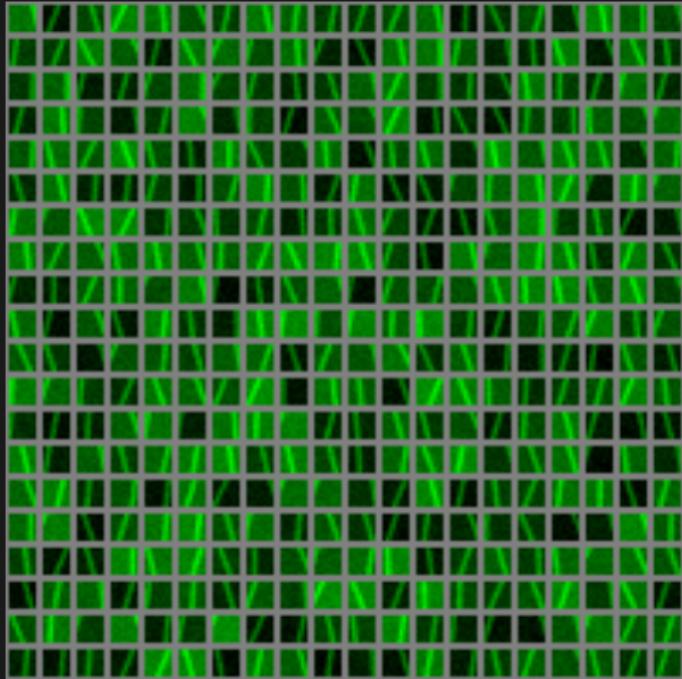


line shape prediction

- **fast** run on straight line, **brake** on curve
- **neural network** for line type classification
 - DenseNet - densely connected convolutional neural network
- **input** 8x8 matrix raw data from line sensors
 - 8 past line positions from 8 sensors
- **output** 5 curves types

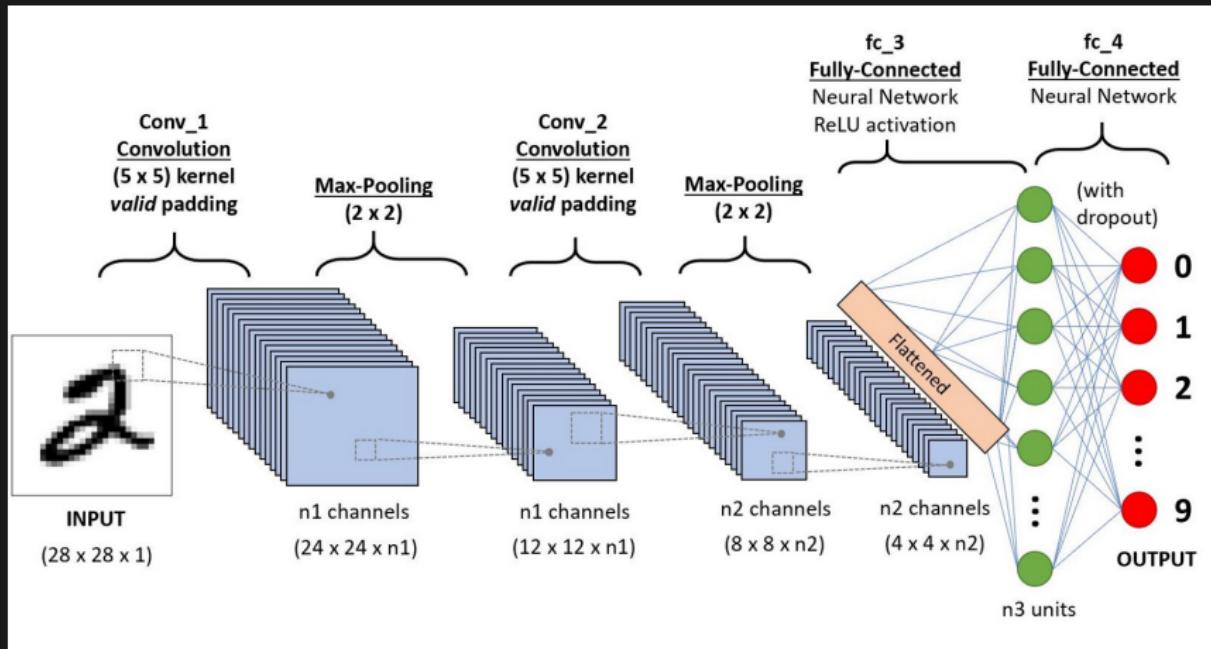


line dataset

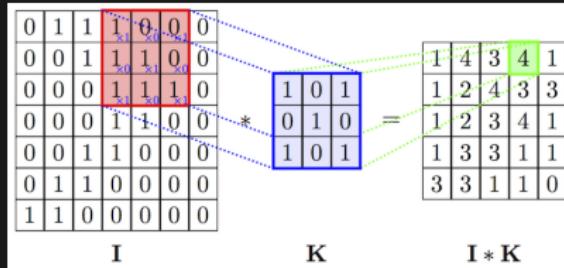


- 20000 for training
- 2500 for testing
- 8x8 inputs
- 5 outputs (5 curves types)
 - two left
 - one straight
 - two right
- augmentation - luma noise, white noise

convolutional neural network - CNN



discrete convolution



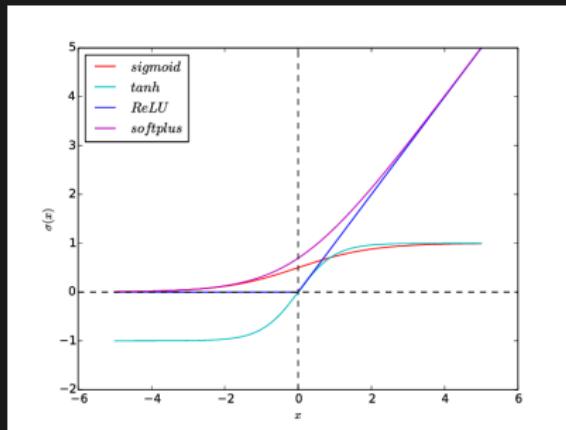
```
for (unsigned y = 0; y < input_height; y++)
for (unsigned x = 0; x < input_width; x++)
{
    float sum = 0.0;

    for (unsigned ky = 0; ky < kernel_height; ky++)
    for (unsigned kx = 0; kx < kernel_width; kx++)
    {
        sum+= kernel[ky][kx]*input[y + ky][x + kx];
    }

    output[y + kernel_height/2][x + kernel_width/2] = sum;
}
```

activation function

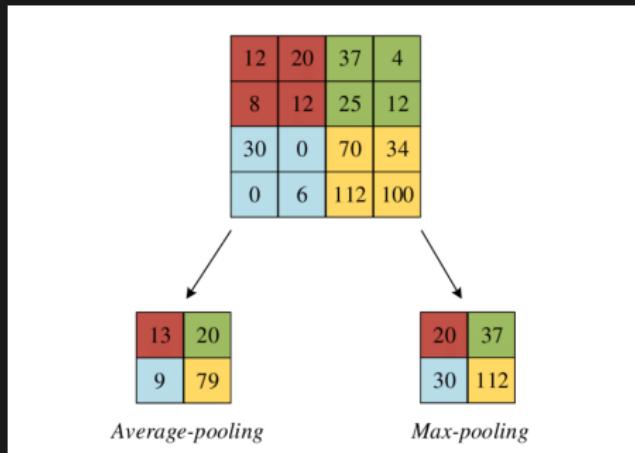
tanh, sigmoid, gaussian, **ReLU**, leaky RELU ...



$$f(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$

$$\frac{df(x)}{dx} = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$

pooling



target hardware

| target | bits | features | frequency |
|----------------------|------|----------------------------------|-----------------|
| AVR Atmega 328 | 8 | single cycle ADD, MUL | 20MHz |
| ARM Cortex M0 | 32 | single cycle ADD, MUL | 48MHz |
| ARM Cortex M4, M7 | 32 | SIMD DUAL 16bit MAC , FPU ... | 72MHz 216MHz |

embedded network implementation

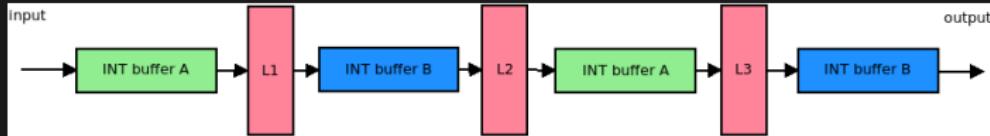
Quantization : convert float weights to `int8_t`

$$scale = \max(|\vec{w}|_1)$$

$$\vec{w}' = \vec{w} \frac{127}{scale}$$

Memory saving : use double buffer memory trick

- `unsigned buffer_size = maxi(layers[i].input_size());`
- `buffer_a = new int8_t(buffer_size);`
- `buffer_b = new int8_t(buffer_size);`



optimize kernel - templates

```
template<unsigned int kernel_size>
void convolution()
{
    for (unsigned y = 0; y < input_height; y++)
        for (unsigned x = 0; x < input_width; x++)
    {
        int sum = 0;

        for (unsigned ky = 0; ky < kernel_size; ky++)
            for (unsigned kx = 0; kx < kernel_size; kx++)
        {
            sum+= kernel[ky][kx]*input[y + ky][x + kx];
        }

        output[y + kernel_size/2][x + kernel_size/2] = (sum*scale)/127;
    }
}

...

convolution <1>(); //1x1 kernel
convolution <3>(); //3x3 kernel
convolution <5>(); //5x5 kernel
```

optimize kernel - unrolling

```
template<unsigned int kernel_size>
void convolution()
{
    for (unsigned y = 0; y < input_height; y++)
        for (unsigned x = 0; x < input_width; x++)
    {
        int sum = 0;

        if (kernel_size == 3)
        {
            sum+= kernel[0][0]*input[y + 0][x + 0];
            sum+= kernel[0][1]*input[y + 0][x + 1];
            sum+= kernel[0][2]*input[y + 0][x + 2];

            sum+= kernel[1][0]*input[y + 1][x + 0];
            sum+= kernel[1][1]*input[y + 1][x + 1];
            sum+= kernel[1][2]*input[y + 1][x + 2];

            sum+= kernel[2][0]*input[y + 2][x + 0];
            sum+= kernel[2][1]*input[y + 2][x + 1];
            sum+= kernel[2][2]*input[y + 2][x + 2];
        }

        output[y + kernel_size/2][x + kernel_size/2] = (sum*scale)/127;
    }
}
```

optimize kernel - SIMD instructions

elementary row MAC operation in 3x3 convolution :

- **sum**+= **kernel**[0][0]***input**[y + 0][x + 0];
- **sum**+= **kernel**[0][1]***input**[y + 0][x + 1];
- **sum**+= **kernel**[0][2]***input**[y + 0][x + 2];

instruction : SMLADX(int32_t val1, int32_t val2, int32_t val3)

- p1 = **val1**[15:0]***val2**[15:0]
- p2 = **val1**[31:16]***val2**[31:16]
- **res**[31:0] = p1 + p2 + **val3**[31:0]

speed up summary

| convolutional | input size | kernels count | kernel size | computing time [ms] | speed up ratio |
|----------------------|------------|---------------|-------------|---------------------|----------------|
| naive | 16x16x16 | 16 | 1x1 | 109.8 | 1.00 |
| naive | 16x16x16 | 16 | 3x3 | 312 | 1.00 |
| naive | 16x16x16 | 16 | 5x5 | 532 | 1.00 |
| template | 16x16x16 | 16 | 1x1 | 26.6 | 4.13 |
| template | 16x16x16 | 16 | 3x3 | 213.6 | 1.46 |
| template | 16x16x16 | 16 | 5x5 | 373.3 | 1.43 |
| template + unrolling | 16x16x16 | 16 | 1x1 | 21 | 5.23 |
| template + unrolling | 16x16x16 | 16 | 3x3 | 57.1 | 5.46 |
| template + unrolling | 16x16x16 | 16 | 5x5 | 106.2 | 5.01 |

network implementation

architecture : IN8x8x1 - C3x3x4 - P2x2 - C3x3x8 - P2x2 - FC5
confussion matrix :

| | | | | |
|---------------|---------------|---------------|---------------|---------------|
| 1029 | 12 | 0 | 0 | 0 |
| 14 | 911 | 18 | 0 | 0 |
| 0 | 21 | 959 | 6 | 0 |
| 0 | 0 | 21 | 968 | 18 |
| 0 | 0 | 0 | 11 | 1012 |
| | | | | |
| 98.658 | 96.504 | 96.092 | 98.274 | 98.252 |

success count 4879

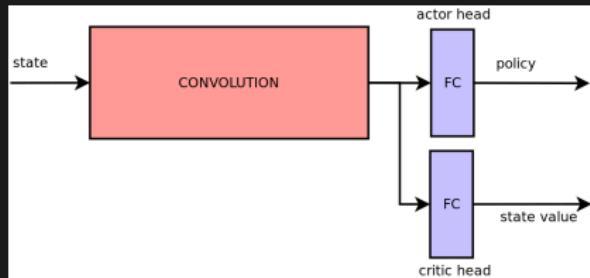
miss count 121

accuracy 97.58%

what is next?

reinforcement learning to take full control -

Advantage Actor Critic - A2C

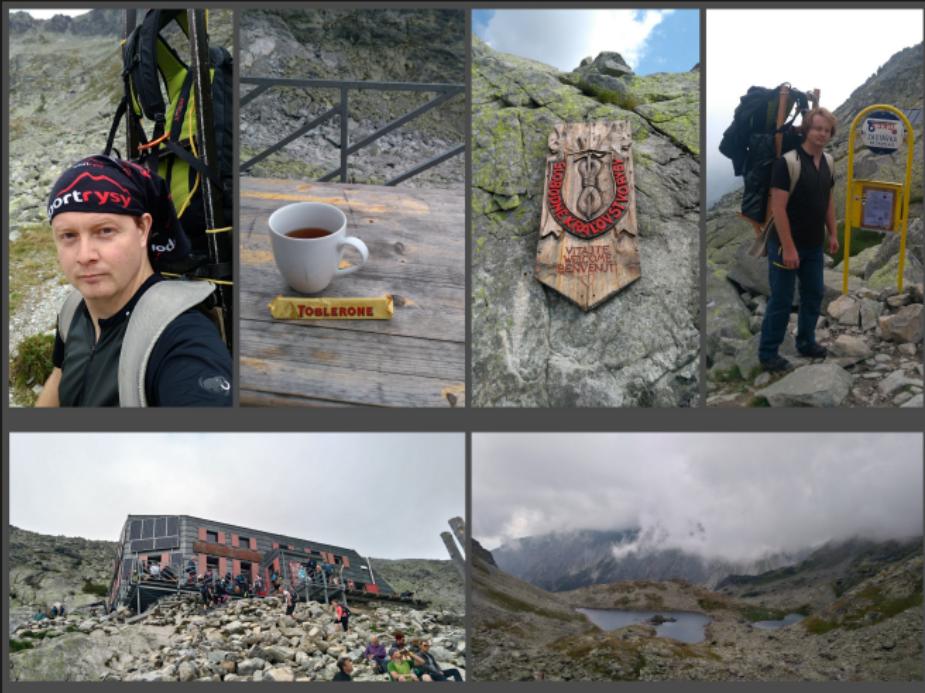


$$\mathcal{L}_{actor} = -(R(n) + \gamma Q(s(n+1); \phi) - Q(s(n); \phi)) \log \pi(s, a; \theta)$$

$$\mathcal{L}_{critic} = (R(n) + \gamma Q(s(n+1); \phi) - Q(s(n); \phi))^2$$

$$\mathcal{L}_{entropy} = \sum_a \pi(s, a; \theta) \log \pi(s, a; \theta)$$

Q&A



michal.nand@gmail.com, github
<https://github.com/michalnand>